

```
In [53]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [54]: df = pd.read_csv('1553768847-housing.csv')
```

```
In [55]: #How many columns
len(df.columns)
```

```
Out[55]: 10
```

```
In [56]: #view first 5 lines of df
df.head(5)
```

```
Out[56]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	hou
0	-122.23	37.88	41	880	129.0	322	
1	-122.22	37.86	21	7099	1106.0	2401	
2	-122.24	37.85	52	1467	190.0	496	
3	-122.25	37.85	52	1274	235.0	558	
4	-122.25	37.85	52	1627	280.0	565	

```
In [57]: #are there duplacate row
sum(df.duplicated())
```

```
Out[57]: 0
```

```
In [58]: # how many null values? Where do the null values exist? Do the empty values affect
df.isnull().sum()
```

```
Out[58]: longitude      0
latitude      0
housing_median_age    0
total_rooms      0
total_bedrooms    207
population      0
households      0
median_income     0
ocean_proximity    0
median_house_value  0
dtype: int64
```

```
In [59]: #how many rows? How many columns?
df.shape
```

```
Out[59]: (20640, 10)
```

```
In [60]: #using mean, median and mode. fill null values givin these data condidions:
#mean - if data is normal distribution and there are few outliers
#median - use if your data is skewed and has outliers
#mode - use when working with categorical data
```

```
In [61]: #lets take a look at the mean value
df['total_bedrooms'].mean()
```

```
Out[61]: 537.8705525375618
```

```
In [62]: #our null values are all in total bedrooms. our data is numerical, so mode is off t
#seeing that houses can have any number of bedrooms(i.e potentially many outliers),
df['total_bedrooms'].median()
```

```
Out[62]: 435.0
```

```
In [63]: #fill in missing values using median
df['total_bedrooms'] = df['total_bedrooms'].fillna(df['total_bedrooms'].median())
```

```
In [64]: #no more missing values, good!
df.isnull().sum()
```

```
Out[64]: longitude          0
latitude          0
housing_median_age  0
total_rooms       0
total_bedrooms    0
population        0
households        0
median_income     0
ocean_proximity   0
median_house_value 0
dtype: int64
```

```
In [65]: # identify unique non-digit values. Determne best way to encode (i.e. turn string v
df['ocean_proximity'].value_counts()
```

```
Out[65]: ocean_proximity
<1H OCEAN    9136
INLAND       6551
NEAR OCEAN   2658
NEAR BAY     2290
ISLAND        5
Name: count, dtype: int64
```

```
In [79]: # replace values on ocean prox with numerical digits
df["ocean_proximity"].replace(['NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'IS
```

```
In [80]: df.head(5)
```

Out[80]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	hou
0	-122.23	37.88	41	880	129.0	322	
1	-122.22	37.86	21	7099	1106.0	2401	
2	-122.24	37.85	52	1467	190.0	496	
3	-122.25	37.85	52	1274	235.0	558	
4	-122.25	37.85	52	1627	280.0	565	

```
In [81]: # x data with house value dropped
x = df.drop('median_house_value',axis=1)

# y data with housing prices
y = df['median_house_value']
```

```
In [82]: x.columns
```

```
Out[82]: Index(['longitude', 'latitude', 'housing_median_age', 'total_rooms',
               'total_bedrooms', 'population', 'households', 'median_income',
               'ocean_proximity'],
              dtype='object')
```

```
In [100... #import train test split, use 20 - 30 % of data in test size
from sklearn.model_selection import train_test_split

x_train , x_test , y_train , y_test = train_test_split(x , y , test_size=0.25 )
```

```
In [101... from sklearn import linear_model
```

```
In [102... #fit training data. keep training and test data seperate, we dont want to give any
reg = linear_model.LinearRegression()
reg.fit(x_train , y_train)
```

```
Out[102... LinearRegression
LinearRegression()
```

```
In [103... y_pred = reg.predict(x_test)
```

```
In [104... y_pred
```

```
Out[104... array([ 77175.32791787, 175177.19766364, 318518.21504612, ...,
        257381.44185985, 31229.57102564, 127814.59296108])
```

```
In [105... #r2 score. want to be close to 1
reg.score(x_test , y_test)
```

```
Out[105... 0.6315061947601698
```

```
In [106... reg.score(x_train , y_train)
```

```
Out[106... 0.637186977923716
```

```
In [107... #use of built in r2 score. get MSE from train and test data  
from sklearn.metrics import mean_absolute_error, r2_score  
  
r2_score(y_test, y_pred)
```

```
Out[107... 0.6315061947601698
```

```
In [108... mean_absolute_error(y_test, y_pred)
```

```
Out[108... 51088.66228130789
```

```
In [109... y_pred_train = reg.predict(x_train)
```

```
In [110... mean_absolute_error(y_train, y_pred_train)
```

```
Out[110... 50982.53564532673
```

```
In [ ]:
```